CASE 2:

Determine Income based on characteristics

- Data Exploration
- Data Splitting
- Data Scaling
- Feature Reduction
- Model Selection
- Parameter Tuning
- Evaluation

CASE 2: Data Exploration

- Census Income Data Set
- https://archive.ics.uci.edu/ml/datasets/Census+Income

CASE 2: Data Exploration

- 15 Attributes:
- age: continuous.
- workclass: Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
- fnlwgt: continuous.
- education: Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
- education-num: continuous.
- marital-status: Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
- occupation: Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
- relationship: Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
- race: White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
- sex: Female, Male.
- capital-gain: continuous.
- capital-loss: continuous.
- hours-per-week: continuous.
- native-country: United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinadad&Tobago, Peru, Hong, Holand-Netherlands.
- class: >50K, <=50K

CASE 2: Data Exploration

Age

Min	17.
1st Qu	28.
Median	37.
Mean	38.5816
3rd Qu	48.
Max	90.

Workclass

Private	22 696
Self-emp-not-inc	2541
Local-gov	2093
?	1836
State-gov	1298
Self-emp-inc	1116
(Other)	981

Education

HS-grad	10 501
Some-college	7291
Bachelors	5355
Masters	1723
Assoc-voc	1382
11th	1175
(Other)	5134

Education-Num

Min	1.
1st Qu	9.
Median	10.
Mean	10.0807
3rd Qu	12.
Max	16.

Marital-Status

Married-civ-spouse	14 976
Never-married	10 683
Divorced	4443
Separated	1025
Widowed	993
Married-spouse-absent	418
Married-AF-spouse	23

Occupation

Prof-specialty	4140
Craft-repair	4099
Exec-managerial	4066
Adm-clerical	3770
Sales	3650
Other-service	3295
(Other)	9541

Relationship

Husband	13 193
Not-in-family	8305
Own-child	5068
Unmarried	3446
Wife	1568
Other-relative	981

Race

White	27 816
Black	3124
Asian-Pac-Islander	1039
Amer-Indian-Eskimo	311
Other	271

Sex

White	27 816
Black	3124
Asian-Pac-Islander	1039
Amer-Indian-Eskimo	311
Other	271

Capital-Gain

1st Qu	0.
3rd Qu	0.
Median	0.
Min	0.
Mean	1077.65
Max	99 999.

Capital-

1st Qu	0.
3rd Qu	0.
Median	0.
Min	0.
Mean	87.3038
Max	4356.

Hours-Per-

Min	1.
1st Qu	40.
Median	40.
Mean	40.4375
3rd Qu	45.
Max	99.

Native-

United-States	29 170
Mexico	643
?	583
Philippines	198
Germany	137
Canada	121
(Other)	1709

Income

<=50K	24 720
>50K	7841

CASE 2: Data Splitting

- Training Set: 32561
- Test Set: 16281

CASE 2: Data Scaling

- Convert categorical variables to numerical variables: workclass, education, marital-status, occupation, relationship, race, sex, native-country, class.

- def parsePoint(line):

CASE 2: Feature Reduction

- fnlwgt
- capital-gain: continuous.
- capital-loss: continuous.

CASE 2: Model Selecting

- LogisticRegressionWithSGD
- Decision Tree Classification
- Bayes

CASE 2: Evaluation

- LogisticRegressionWithSGD

```
Nings-MBP:bin ningzhang$ ./spark-submit ../../../../../Users/ningzhang/Desktop/midterm/LogisticRegressionWithSG
32561
16281
Training Error = 0.18051716725
```

- Decision Tree Classification

```
Nings-MBP:bin ningzhang$ ./spark-submit ../../../../../Users/ningzhang/Desktop/midterm/DecitionTreeClaasificatio
32561
16281
Test Error = 0.827160493827
Learned classification tree model:
DecisionTreeModel classifier of depth 10 with 1199 nodes
  If (feature 3 <= 12.0)
   If (feature 0 <= 33.0)
    If (feature 0 <= 26.0)
     If (feature 9 <= 43.0)
      If (feature 0 <= 24.0)
       If (feature 0 <= 21.0)
        If (feature 9 <= 39.0)
         Predict: 1.0
```

- Bayes

```
Nings-MBP:bin ningzhang$ ./spark-submit ../../../../../Users/ningzhang/Desktop/midterm/Bayes.
Train Data:
32561
Test Data:
16281
Accuracy: 0.117928874148
```

CASE 2: Evaluation